

lab08 New System Call

软件工程 2018 级 1813075 刘茵

一、实验目标

- Add a new system call with arguments into the linux kernel
- The new system call will return all processes information to user mode

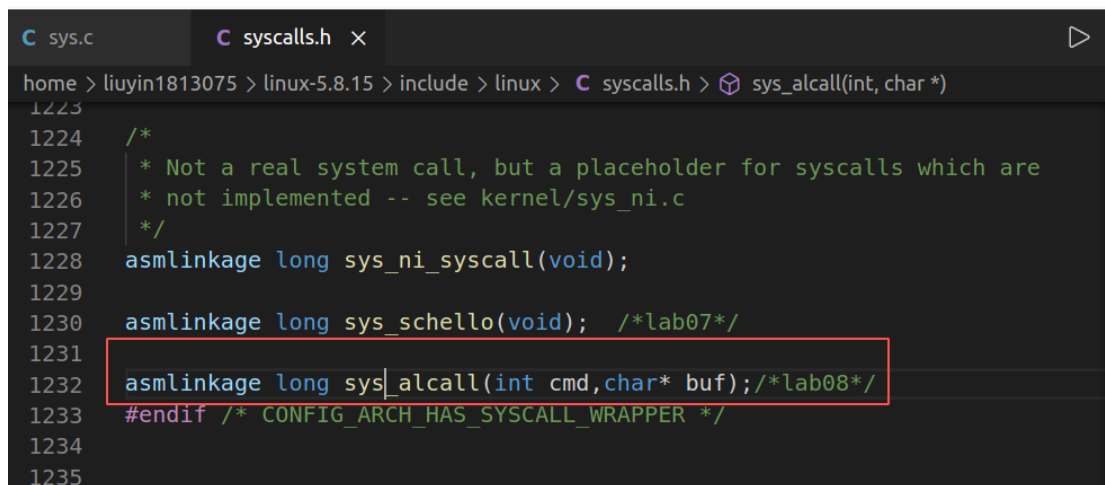
二、操作过程

(一) PART 1

1. 进入 USR/SRC/LINUX 目录 (和 HOME/LINUX5.8.15 相同)

```
liuyin1813075@liuyin-VirtualBox:~/linux-5.8.15$ cd include/linux/  
liuyin1813075@liuyin-VirtualBox:~/linux-5.8.15/include/linux$ code syscalls.h
```

2. 用 vscode 打开 include/linux/syscall.h 文件, 并修改



```
C sys.c  C syscalls.h x  
home > liuyin1813075 > linux-5.8.15 > include > linux > C syscalls.h > sys_alcall(int, char *)  
1223  
1224 /*  
1225  * Not a real system call, but a placeholder for syscalls which are  
1226  * not implemented -- see kernel/sys_ni.c  
1227  */  
1228 asmlinkage long sys_ni_syscall(void);  
1229  
1230 asmlinkage long sys_schello(void); /*lab07*/  
1231  
1232 asmlinkage long sys_alcall(int cmd, char* buf); /*lab08*/  
1233 #endif /* CONFIG_ARCH_HAS_SYSCALL_WRAPPER */  
1234  
1235
```

3. 修改/kernel/sys.c 文件, 添加 SYSCALL_DEFINE3 (alcall,int,cmd,char*,buf)函数

```
liuyin1813075@liuyin-VirtualBox:~/linux-5.8.15$ cd kernel/  
liuyin1813075@liuyin-VirtualBox:~/linux-5.8.15/kernel$ code sys.c
```

构造了一个结构体, 其中含有一个 buf2 数组和代表数组中进程数的 int 型变量 num;

利用 copy_to_user 和 put_user 两个函数实现内核空间数据与用户空间数据的相互访问。

```

/*lab08 new syscall*/
struct process{
    char buf2[1024];
    int num;
}
SYSCALL_DEFINE2(alcall, int, cmd, char *, buf)
{
    char temp[256];
    struct process allps;
    struct task_struct *p;
    snprintf(temp,256, "Hello new system call alcall (%d, %x)!\n", cmd, buf);
    printk("%s\n",temp);
    strcat(allps.buf2,temp);
    printk("LiuYin 1813075\n");
    snprintf(temp,256, "%-20s %-6s %-6s\n", "Name", "Pid", "Stat");
    printk("%s\n",temp);
    strcat(allps.buf2,temp);
    int count=0;
    for (p = &init_task; (p = next_task(p)) != &init_task;) {
        snprintf(temp, sizeof(temp), "%-20s %-6d %-6ld\n", p->comm,
            p->pid, p->state);
        printk("%s\n",temp);
        strcat(allps.buf2, temp);
        count++;
        if (count >= cmd)
            break;
    }
    allps.num=count;
    int ret=0;
    printk("Ready to copy to user");
    ret=copy_to_user(((struct process*)buf)->buf2,allps.buf2,strlen(allps.buf2));
    put_user(count,&(((struct process*)buf)->num));
    return ret;
}

```

4. 修改 arch/x86/entry/syscalls/syscall_64.tbl 文件，增加 441

```

liuyin1813075@liuyin-VirtualBox:~/linux-5.8.15$ cd arch/x86/entry/syscalls/
liuyin1813075@liuyin-VirtualBox:~/linux-5.8.15/arch/x86/entry/syscalls$ ls
arch Makefile syscall_32.tbl syscall_64.tbl syscallhdr.sh syscalltbl.sh
liuyin1813075@liuyin-VirtualBox:~/linux-5.8.15/arch/x86/entry/syscalls$ code sys
call_64.tbl

```

syscall_64.tbl - Visual Studio Code				
File Edit Selection View Go Run Terminal Help				
home > liuyin1813075 > linux-5.8.15 > arch > x86 > entry > syscalls > syscall_64.tbl				
355	431	common	fsconfig	sys_fsconfig
356	432	common	fsmount	sys_fsmount
357	433	common	fspick	sys_fspick
358	434	common	pidfd_open	sys_pidfd_open
359	435	common	clone3	sys_clone3
360	437	common	openat2	sys_openat2
361	438	common	pidfd_getfd	sys_pidfd_getfd
362	439	common	faccessat2	sys_faccessat2
363				
364	440	common	schello	sys_schello
365	441	common	alcall	__x64_sys_alcall
366				#

5. re-configure the kernel

```
cp linux_module .config  
make oldconfig  
make gconfig
```

6. 依次运行以下命令编译内核。

```
make clean  
make -j5  
sudo make modules_install  
sudo make install
```

```
OBJCOPY arch/x86/boot/setup.bin  
BUILD arch/x86/boot/bzImage  
Setup is 14140 bytes (padded to 14336 bytes).  
System is 8937 kB  
CRC 374401d9  
Kernel: arch/x86/boot/bzImage is ready (#5)
```

```
Kernel: arch/x86/boot/bzImage is ready (#5)  
liuyin1813075@liuyin-VirtualBox:~/linux-5.8.15$ sudo make modules_install  
[sudo] liuyin1813075 的密码:  
INSTALL drivers/thermal/intel/x86_pkg_temp_thermal.ko  
INSTALL fs/efivarfs/efivarfs.ko  
INSTALL net/ipv4/netfilter/iptables_nat.ko  
INSTALL net/ipv4/netfilter/nf_log_arp.ko  
INSTALL net/ipv4/netfilter/nf_log_ipv4.ko  
INSTALL net/ipv6/netfilter/nf_log_ipv6.ko  
INSTALL net/netfilter/nf_log_common.ko  
INSTALL net/netfilter/xt_LOG.ko  
INSTALL net/netfilter/xt_MASQUERADE.ko  
INSTALL net/netfilter/xt_addrtype.ko  
INSTALL net/netfilter/xt_mark.ko  
INSTALL net/netfilter/xt_nat.ko  
DEPMOD 5.8.15Lee202009
```

```
liuyin1813075@liuyin-VirtualBox:~/linux-5.8.15$ sudo make install  
sh ./arch/x86/boot/install.sh 5.8.15Lee202009 arch/x86/boot/bzImage \  
System.map "/boot"  
run-parts: executing /etc/kernel/postinst.d/apt-auto-removal 5.8.15Lee202009 /boot/vmlinuz-5.8.15Lee202009  
run-parts: executing /etc/kernel/postinst.d/initramfs-tools 5.8.15Lee202009 /boot/vmlinuz-5.8.15Lee202009  
update-initramfs: Generating /boot/initrd.img-5.8.15Lee202009  
find: '/var/tmp/mkinitramfs_eUEU4S/lib/modules/5.8.15Lee202009/kernel': 没有那个文件或目录  
run-parts: executing /etc/kernel/postinst.d/unattended-upgrades 5.8.15Lee202009 /boot/vmlinuz-5.8.15Lee202009  
run-parts: executing /etc/kernel/postinst.d/update-notifier 5.8.15Lee202009 /boot/vmlinuz-5.8.15Lee202009  
run-parts: executing /etc/kernel/postinst.d/vboxadd 5.8.15Lee202009 /boot/vmlinuz-5.8.15Lee202009  
VirtualBox Guest Additions: Building the modules for kernel 5.8.15Lee202009.  
  
VirtualBox Guest Additions: Look at /var/log/vboxadd-setup.log to find out what went wrong  
run-parts: executing /etc/kernel/postinst.d/zz-update-grub 5.8.15Lee202009 /boot/vmlinuz-5.8.15Lee202009  
Sourcing file '/etc/default/grub'  
Sourcing file '/etc/default/grub.d/init-select.cfg'  
正在生成 grub 配置文件 ...  
找到 Linux 镜像: /boot/vmlinuz-5.8.15Lee202009  
找到 initrd 镜像: /boot/initrd.img-5.8.15Lee202009  
找到 Linux 镜像: /boot/vmlinuz-5.8.15  
找到 initrd 镜像: /boot/initrd.img-5.8.15  
找到 Linux 镜像: /boot/vmlinuz-5.8.15.old  
找到 initrd 镜像: /boot/initrd.img-5.8.15  
找到 Linux 镜像: /boot/vmlinuz-5.4.0-54-generic  
找到 initrd 镜像: /boot/initrd.img-5.4.0-54-generic  
找到 Linux 镜像: /boot/vmlinuz-5.4.0-53-generic  
找到 initrd 镜像: /boot/initrd.img-5.4.0-53-generic  
找到 Linux 镜像: /boot/vmlinuz-5.4.0-42-generic  
找到 initrd 镜像: /boot/initrd.img-5.4.0-42-generic  
Found memtest86+ image: /boot/memtest86+.elf  
Found memtest86+ image: /boot/memtest86+.bin  
完成
```

7. 编写 testalcall.c 并运行。

```
home > liuyin1813075 > oscourse > course8 > C testalcall.c > main(int, char *[])
1  #include<unistd.h>
2  #include<sys/syscall.h>
3  #include<sys/types.h>
4  #include<stdio.h>
5  #define __NR_alcall 438
6  struct process{
7      char buf2[1025];
8      int num;
9  }
10 long alcall(int cmd,char *buf){
11     return syscall(__NR_alcall,cmd,buf);
12 }
13 int main(int argc,char *argv[]){
14     struct process result;
15     int cmd;
16     cmd=9;
17     alcall(cmd,&result);
18     printf("ok!run dmesg | grep alcall in terminal!%s\n",result.buf2);
19     printf("the number of processes is %d\n",result.num);
20     return 0;
21 }
```

编译并输出结果。

```
liuyin1813075@liuyin-VirtualBox:~/oscourse/course8$ ./testalcall
ok!run dmesg | grep alcall in terminal! Hello new system call alcall (9, 6b0f152
0)!
Name                Pid    Stat
systemd              1      1
kthreadd             2      1
rcu_gp               3     1026
rcu_par_gp           4     1026
kworker/0:0H         6     1026
kworker/u8:0         7     1026
mm_percpu_wq         8     1026
ksoftirqd/0          9      1
rcu_sched            10     1026

the number of processes is 32765
```

系统调用输出。

```
liuyin1813075@liuyin-VirtualBox:~/oscourse/course8$ dmesg | grep alcall
[ 99.101930] Hello new system call alcall (9, dcf16ef0)!
liuyin1813075@liuyin-VirtualBox:~/oscourse/course8$ dmesg
[ 0.000000] Linux version 5.8.15Lee202009 (liuyin1813075@liuyin-VirtualBox) (
gcc (Ubuntu 9.3.0-17ubuntu1~20.04) 9.3.0, GNU ld (GNU Binutils for Ubuntu) 2.34)
#5 SMP Fri Dec 4 16:59:52 CST 2020
```

[2028.626189] Hello new system call alcall (9, 3ae59b30)!

[2028.626193] LiuYin 1813075

[2028.626198]	Name	Pid	Stat
----------------	------	-----	------

[2028.626202]	systemd	1	1
----------------	---------	---	---

[2028.626207]	kthreadd	2	1
----------------	----------	---	---

[2028.626211]	rcu_gp	3	1026
----------------	--------	---	------

[2028.626215]	rcu_par_gp	4	1026
----------------	------------	---	------

[2028.626219]	kworker/0:0H	6	1026
----------------	--------------	---	------

[2028.626223]	kworker/u8:0	7	1026
----------------	--------------	---	------

[2028.626227]	mm_percpu_wq	8	1026
----------------	--------------	---	------

[2028.626231]	ksoftirqd/0	9	1
----------------	-------------	---	---

[2028.626236]	rcu_sched	10	1026
----------------	-----------	----	------